

Page: [Details]

Result: 474 - multithread_attn_kernel

▼

Add Baseline

Apply Rules

Occupancy Calculator

Source Comparison

Save as PDF

Time

Cycles

GPU

SM Frequency

Process

Attributes

4253 - multithread_attn_kernel (6, 1, 1)x(256, 1, 1)

6.11 usecond

9,848

0 - NVIDIA GeForce RTX 4060 Laptop GPU

1.61 cycle/nsecond

[11435] run

GPU Speed of Light Throughput

GPU Throughput Rooflines

High-level overview of the throughput for compute and memory resources of the GPU. For each unit, the throughput reports the achieved percentage of utilization with respect to the theoretical maximum. Breakdowns show the throughput for each individual sub-metric of Compute and Memory to clearly identify the highest contributor. High-level overview of the utilization for compute and memory resources of the GPU presented as a roofline chart.

Compute (SM) Throughput [%]

Memory Throughput [%]

L1/TEX Cache Throughput [%]

L2 Cache Throughput [%]

DRAM Throughput [%]

2.82

2.82

14.09

0.76

1.52

Duration [usecond]

Elapsed Cycles [cycle]

SM Active Cycles [cycle]

SM Frequency [cycle/nsecond]

DRAM Frequency [cycle/nsecond]

6.11

9,848

1,969.62

1.61

6.95

Small Grid

This kernel grid is too small to fill the available resources on this device, resulting in only 0.1 full waves across all SMs. Look at [Launch Statistics](#) for more details.

Roofline Analysis

The ratio of peak float (fp32) to double (fp64) performance on this device is 64:1. The kernel achieved close to 0% of this device's fp32 peak performance and 0% of its fp64 peak performance. See the [Kernel Profiling Guide](#) for more details on roofline analysis.

Floating Point Operations Roofline

Performance [FLOP/s] (1 = 100,000,000,000)

Arithmetic Intensity [FLOP/byte]

PM Sampling

Timeline view of PM metrics sampled periodically over the workload duration. Data is collected across multiple passes. Use this section to understand how workload behavior changes over its runtime.

Maximum Sampling Interval [usecond]

Maximum Buffer Size [Mbytes]

1

4

Pass Groups

Dropped Samples [sample]

2

0

Compute Workload Analysis

Detailed analysis of the compute resources of the streaming multiprocessors (SM), including the achieved instructions per clock (IPC) and the utilization of each available pipeline. Pipelines with very high utilization might limit the overall performance.

Executed Ipc Elapsed [inst/cycle]

Executed Ipc Active [inst/cycle]

Issued Ipc Active [inst/cycle]

0.06

0.32

0.33

SM Busy [%]

Issue Slots Busy [%]

8.24

8.24

Low Utilization

Est. Local Speedup: 94.10%

All compute pipelines are under-utilized. Either this kernel is very small or it doesn't issue enough warps per scheduler. Check the [Launch Statistics](#) and [Scheduler Statistics](#) sections for further details.

Pipe Utilization (% of active cycles)

Pipe Utilization (% of peak instructions executed)

ALU

FMA

FP64

Tensor (All)

Tensor (FP)

Tensor (INT)

0.0

25.0

50.0

75.0

100.0

Utilization [%]

Utilization [%]

Memory Workload Analysis

Detailed analysis of the memory resources of the GPU. Memory can become a limiting factor for the overall kernel performance when fully utilizing the involved hardware units (Mem Busy), exhausting the available communication bandwidth between those units (Max Bandwidth), or by reaching the maximum throughput of issuing memory instructions (Mem Pipes Busy). Detailed chart of the memory units. Detailed tables with data for each memory unit.

Memory Throughput [Gbyte/second]

L1/TEX Hit Rate [%]

L2 Hit Rate [%]

L2 Compression Success Rate [%]

3.39

73.88

33.43

0

Mem Busy [%]

Max Bandwidth [%]

Mem Pipes Busy [%]

L2 Compression Ratio

1.82

2.82

2.82

0

L1TEX Global Load Access Pattern

Est. Speedup: 2.25%

The memory access pattern for global loads from L1TEX might not be optimal. On average, only 6.4 of the 32 bytes transmitted per sector are utilized by each thread. This could possibly be caused by a stride between threads. Check the [Source Counters](#) section for uncoalesced global loads.

L1TEX Global Store Access Pattern

Est. Speedup: 0.06%

The memory access pattern for global stores to L1TEX might not be optimal. On average, only 31.3 of the 32 bytes transmitted per sector are utilized by each thread. This could possibly be caused by a stride between threads. Check the [Source Counters](#) section for uncoalesced global stores.

Shared Memory

	Instructions	Requests	Wavefronts	% Peak	Bank Conflicts
Shared Load	384	384			
Shared Load Matrix	0	0	384	0.16	0
Shared Store	240	240			
Shared Store From Global Load	0	0	0	0	0
Shared Atomic	0	0	0	0	0
Other	-	-	2,052	0.94	0
Total	624	624	2,676	1.13	0

L1/TEX Cache

	Instructions	Requests	Wavefronts	% Peak	Sectors	Sectors/Req	Hit Rate	Bytes	Sector Misses to L2	% Peak to L2	Returns to SM
Local Load	0	0	0	0	0	0	0	0			
Global Load	666	666			1,314	1.97	80.82	42,048	252	0.11	
Global Load To Shared Store (access)	0	0	795	0.34	0	0		0			
Global Load To Shared Store (bypass)	0	0			0	0	-	0	0	0	
Surface Load	0	0	0	0	0	0	0	0	0	0	
Texture Load	0	0	0	0	0	0	0	0	0	0	
Global Store	120	120	120	0.05	432	3.60	52.78	13,824			
Local Store	0	0	0	0	0	0	0	0	432	0.18	
Surface Store	0	0	0	0	0	0	0	0	0	0	
Global Reduction	0	0	0	0	0	0	0	0	0	0	
DSMEM Reduction	0	0	0	0	0	-	-	0	0	0	
Surface Reduction	0	0	0	0	0	0	0	0	0	0	
Global Atomic ALU	0	0	0	0	0	0	0	0	0	0	
Global Atomic CAS	0	0	0	0	0	0	0	0	0	0	see r
Surface Atomic ALU	0	0	0	0	0	0	0	0	0	0	see r
Surface Atomic CAS	0	0	0	0	0	0	0	0	0	0	
Loads	666	666	795	0.34	1,314	1.97	80.82	42,048	252	0.11	
Stores	120	120	120	0.05	432	3.60	52.78	13,824	432	0.18	
Stores & Reductions	0	0	0	0	0	0	0	0	0	0	

L2 Cache

	Requests	Sectors	Sectors/Req	% Peak	Hit Rate	Bytes	Throughput	Sector Misses to Device	Sector Misses to System	Sector Misses to Peer
L1/TEX Load	189	252	1.33	0.08	0	8,064	1,319,371,727.75	252	0	0
L1/TEX Store	121	432	3.57	0.28	100	13,824	2,261,780,104.71	0	0	0
L1/TEX Atomic ALU	0	0	0	0	0	0	0	0	0	0
L1/TEX Atomic CAS	0	0	0	0	0	0	0	0	0	0
L1/TEX Reduction	0	0	0	0	0	0	0	0	0	0
L1/TEX Total	311	684	2.20	0.22	63.16	21,888	3,581,151,832.46	252	0	0
ECC Total	-	0	-	0	-	0	0	0	-	-
GPU Total	502	2,342	4.67	0.76	75.43	74,944	12,261,780,104.71	255	0	0

L2 Cache Eviction Policies

	First	Hit Rate	Last	Hit Rate	Normal	Hit Rate	Normal Demote	Hit Rate
L1/TEX Load	0	0	0	0	252	0	0	0
L1/TEX Store	0	0	0	0	432	100	0	0
L1/TEX Atomic	0	0	0	0	0	0	-	-
L1/TEX Total	0	0	0	0	684	63.16	0	0
GPU Total	176	84.09	0	0	873	70.79	0	0

Device Memory

	Sectors	% Peak	Bytes	Throughput
Load	648	1.52	20,736	3,392,670,157.07
Store	0	0	0	0
Total	648	1.52	20,736	3,392,670,157.07

Scheduler Statistics

Summary of the activity of the schedulers issuing instructions. Each scheduler maintains a pool of warps that it can issue instructions for. The upper bound of warps in the pool (Theoretical Warps) is limited by the launch configuration. On every cycle each scheduler checks the state of the allocated warps in the pool (Active Warps). Active warps that are not stalled (Eligible Warps) are ready to issue their next instruction. From the set of eligible warps the scheduler selects a single warp from which to issue one or more instructions (Issued Warp). On cycles with no eligible warps, the issue slot is skipped and no instruction is issued. Having many skipped issue slots indicates poor latency hiding.

Active Warps Per Scheduler [warp]

Eligible Warps Per Scheduler [warp]

Issued Warp Per Scheduler

1.91

0.11

0.10

No Eligible [%]

One or More Eligible [%]

90.29

9.71

Issue Slot Utilization

Est. Local Speedup: 90.29%

Every scheduler is capable of issuing one instruction per cycle, but for this kernel each scheduler only issues an instruction every 10.3 cycles. This might leave hardware resources underutilized and may lead to less optimal performance. Out of the maximum of 12 warps per scheduler, this kernel allocates an average of 1.91 active warps per scheduler, but only an average of 0.11 warps were eligible per cycle. Eligible warps are the subset of active warps that are ready to issue their next instruction. Every cycle with no eligible warp results in no instruction being issued and the issue slot remains unused. To increase the number of eligible warps, reduce the time the active warps are stalled by inspecting the top stall reasons on the [Warp State Statistics](#) and [Source Counters](#) sections.

Warp State Statistics

Analysis of the states in which all warps spent cycles during the kernel execution. The warp states describe a warp's readiness or inability to issue its next instruction. The warp cycles per instruction define the latency between two consecutive instructions. The higher the value, the more warp parallelism is required to hide this latency. For each warp state, the chart shows the average number of cycles spent in that state per issued instruction. Stalls are not always impacting the overall performance nor are they completely avoidable. Only focus on stall reasons if the schedulers fail to issue every cycle. When executing a kernel with mixed library and user code, these metrics show the combined values.

Warp Cycles Per Issued Instruction [cycle]

Warp Cycles Per Executed Instruction [cycle]

19.72

20.33

Avg. Active Threads Per Warp

Avg. Not Predicated Off Threads Per Warp

27.32

19.32

Barrier Stalls

Est. Speedup: 48.74%

On average, each warp of this kernel spends 9.6 cycles being stalled waiting for sibling warps at a CTA barrier. A high number of warps waiting at a barrier is commonly caused by diverging code paths before a barrier. This causes some warps to wait a long time until other warps reach the synchronization point. Whenever possible, try to divide up the work into blocks of uniform workloads. If the block size is 512 threads or greater, consider splitting it into smaller groups. This can increase eligible warps without affecting occupancy, unless shared memory becomes a new occupancy limiter. Also, try to identify which barrier instruction causes the most stalls, and optimize the code executed before that synchronization point first. This stall type represents about 48.7% of the total average of 19.7 cycles between issuing two instructions.

Warp Stall

Check the [Warp Stall Sampling \(All Samples\)](#) table for the top stall locations in your source based on sampling data. The [Kernel Profiling Guide](#) provides more details on each stall reason.

Thread Divergence

Est. Speedup: 1.12%

Instructions are executed in warps, which are groups of 32 threads. Optimal instruction throughput is achieved if all 32 threads of a warp execute the same instruction. The chosen launch configuration, early thread completion, and divergent flow control can significantly lower the number of active threads in a warp per cycle. This kernel achieves an average of 27.3 threads being active per cycle. This is further reduced to 19.3 threads per warp due to predication. The compiler may use predication to avoid an actual branch. Instead, all instructions are scheduled, but a per-thread condition code or predicate controls which threads execute the instructions. Try to avoid different execution paths within a warp when possible.

Instruction Statistics

Statistics of the executed low-level assembly instructions (SASS). The instruction mix provides insight into the types and frequency of the executed instructions. A narrow mix of instruction types implies a dependency on few instruction pipelines, while others remain unused. Using multiple pipelines allows hiding latencies and enables parallel execution. Note that 'Instructions/Opcode' and 'Executed Instructions' are measured differently and can diverge if cycles are spent in system calls.

Executed Instructions [inst]

Issued Instructions [inst]

15,120

15,588

Avg. Executed Instructions Per Scheduler [inst]

Avg. Issued Instructions Per Scheduler [inst]

157.50

162.38

FP32 Non-Fused Instructions

Est. Speedup: 0.43%

This kernel executes 324 fused and 462 non-fused FP32 instructions. By converting pairs of non-fused instructions to their [Fused](#), higher-throughput equivalent, the achieved FP32 performance could be increased by up to 29% (relative to its current performance). Check the Source page to identify where this kernel executes FP32 instructions.

NVLink Topology

NVLink Topology diagram shows logical NVLink connections with transmit/receive throughput.

NVLink Tables

Detailed tables with properties for each NVLink.

NUMA Affinity

Non-uniform memory access (NUMA) affinities based on compute and memory distances for all GPUs.

Launch Statistics

Summary of the configuration used to launch the kernel. The launch configuration defines the size of the kernel grid, the division of the grid into blocks, and the GPU resources needed to execute the kernel. Choosing an efficient launch configuration maximizes device utilization.

Grid Size

Registers Per Thread [register/thread]

Block Size

Threads [thread]

Waves Per SM

Uses Green Context

6

64

256

1,536

0.06

0

Function Cache Configuration

Static Shared Memory Per Block [byte/block]

Dynamic Shared Memory Per Block [kbyte/block]

Driver Shared Memory Per Block [kbyte/block]

Shared Memory Configuration Size [Kbyte]

SMs [SM]

CachePreferNone

0

1.02

1.02

16.38

24

Small Grid

Est. Speedup: 75.00%

The grid for this launch is configured to execute only 6 blocks, which is less than the GPU's 24 multiprocessors. This can underutilize some multiprocessors. If you do not intend to execute this kernel concurrently with other workloads, consider reducing the block size to have at least one block per multiprocessor or increase the size of the grid to fully utilize the available hardware resources. See the [Hardware Model](#) description for more details on launch configurations.

Occupancy

Occupancy is the ratio of the number of active warps per multiprocessor to the maximum number of possible active warps. Another way to view occupancy is the percentage of the hardware's ability to process warps that is actively in use. Higher occupancy does not always result in higher performance, however, low occupancy always reduces the ability to hide latencies, resulting in overall performance degradation. Large discrepancies between the theoretical and the achieved occupancy during execution typically indicates highly imbalanced workloads.

Theoretical Occupancy [%]

Theoretical Active Warps per SM [warp]

Achieved Occupancy [%]

Achieved Active Warps Per SM [warp]

66.67

32

13.31

6.39

Block Limit Registers [block]

Block Limit Shared Mem [block]

Block Limit Warps [block]

Block Limit SM [block]

4

8

6

24

Achieved Occupancy

Est. Speedup: 80.03%

The difference between calculated theoretical (66.7%) and measured achieved occupancy (13.3%) can be the result of warp scheduling overheads or workload imbalances during the kernel execution. Load imbalances can occur between warps within a block as well as across blocks of the same kernel. See the [CUDA Best Practices Guide](#) for more details on optimizing occupancy.

Theoretical Occupancy

Est. Speedup: 33.33%

The 8.00 theoretical warps per scheduler this kernel can issue according to its occupancy are below the hardware maximum of 12. This kernel's theoretical occupancy (66.7%) is limited by the number of required registers.

GPU and Memory Workload Distribution

Analysis of workload distribution in active cycles of SM, SMP, SMSP, L1 & L2 caches, and DRAM

Average SM Active Cycles [cycle]

Average L2 Active Cycles [cycle]

Average DRAM Active Cycles [cycle]

Total L1 Elapsed Cycles [cycle]

Total SMSP Elapsed Cycles [cycle]

1,969.62

1,497.88

648

236,312

945,248

Average L1 Active Cycles [cycle]

Average SMSP Active Cycles [cycle]

Total SM Elapsed Cycles [cycle]

Total L2 Elapsed Cycles [cycle]

Total DRAM Elapsed Cycles [cycle]

1,969.62

1,672.25

236,210

154,240

169,984

SMs Workload Imbalance

Est. Speedup: 15.07%

One or more SMs have a much lower number of active cycles than the average number of active cycles. Maximum instance value is 75.34% above the average, while the minimum instance value is 100.00% below the average.

SMSPs Workload Imbalance

Est. Speedup: 13.22%

One or more SMSPs have a much lower number of active cycles than the average number of active cycles. Maximum instance value is 77.82% above the average, while the minimum instance value is 100.00% below the average.

L1 Slices Workload Imbalance

Est. Speedup: 15.07%

One or more L1 Slices have a much lower number of active cycles than the average number of active cycles. Maximum instance value is 75.34% above the average, while the minimum instance value is 100.00% below the average.

L2 Slices Workload Imbalance

Est. Speedup: 7.08%

One or more L2 Slices have a much lower number of active cycles than the average number of active cycles. Maximum instance value is 45.55% above the average, while the minimum instance value is 87.78% below the average.

Source Counters

Source metrics, including branch efficiency and sampled warp stall reasons. Warp Stall Sampling metrics are periodically sampled over the kernel runtime. They indicate when warps were stalled and couldn't be scheduled. See the documentation for a description of all stall reasons. Only focus on stalls if the schedulers fail to issue every cycle.

Branch Instructions [inst]

Branch Instructions Ratio [%]

1,800

0.12

Branch Efficiency [%]

Avg. Divergent Branches

98.32

0.25

Uncoalesced Global Accesses

Est. Speedup: 5.13%

This kernel has uncoalesced global accesses resulting in a total of 576 excessive sectors (33% of the total 1746 sectors). Check the L2 Theoretical Sectors Global Excessive table for the primary source locations. The [CUDA Programming Guide](#) has additional information on reducing uncoalesced device memory accesses.

L2 Theoretical Sectors Global Excessive

Location	Value	Value (%)
0x7af63f283ba0 in multithread_attn_kernel	36	6
0x7af63f283b80 in multithread_attn_kernel	36	6
0x7af63f283b60 in multithread_attn_kernel	36	6
0x7af63f283b40 in multithread_attn_kernel	36	6
0x7af63f283b20 in multithread_attn_kernel	36	6

Follow the rules outputs to get guidance on how to navigate through the report and quickly discover performance bottlenecks in this kernel.

You could also disable [Individual Sections](#) to focus on selected performance aspects and make profiling faster.